SafeScreen — Roadmap to Real Earning Potential

🎯 **The decision:** SafeScreen is the project to make real. It's the only one with a recurring B2B market and **zero financial/legal risk** — non-negotiable for your situation. The engine already exists; it needs proof and a product wrapper. That's the work only you can do.

Where it actually stands (diagnosed today)

Component	Status
Trained classifier (`toxic_detector.onnx`)	✅ Present, runs
OCR models (PP-OCRv3)	✅ Present
Inference + speed benchmarks	✅ Present (but messy — 6 duplicate entrypoints)
Accuracy / correctness proof	❌ **None. This is the gap.**
Dataset record	❌ `total_images: 0` — no documented training/eval set
Class taxonomy	⚠️ Inconsistent across files
README / docs	❌ None

Translation: it's fast, but unproven. Speed without proven accuracy = not sellable to a business. Fix that one thing and everything downstream unlocks.

The core principle

> Nobody pays monthly for "a model." They pay for **"I can trust this to catch X% of bad content and not falsely flag good content."** A number they can put in a

> compliance doc. Your entire job is to produce that trustworthy number — then wrap it so it's easy to buy.

Phase 0 — Clean the foundation (1 lucid day)

- [ ] Lock ONE final class taxonomy (pick: `violence`, `gore`, `nsfw` — drop `poop/vomit/medical/biological` unless a buyer needs them)
- [ ] Consolidate to **one** inference entrypoint; delete the `_bak` / duplicate scripts
- [ ] Confirm the ONNX model loads and runs one image end-to-end (install `onnxruntime`)
- [ ] Write a short honest README: what it detects, how to run it, current known limits

Phase 1 — VALIDATION (the make-or-break — 3–5 lucid days)

This is the part with your name on it. An agent can build the harness; only you can judge real-world edge cases and source honest test data.

- [ ] Assemble a **held-out test set** you did NOT train on: ~200–500 labeled images per class + a big "safe" set (safe images are where false positives hurt)
- [ ] Build an eval script that outputs: **precision, recall, F1 per class, a confusion matrix, and the false-positive rate on safe images**
- [ ] Run it. Write the numbers down honestly.
- [ ] **Honest gate:** if precision/recall are strong → you have a product. If weak → that's a finding, not a failure: either retrain with more data or narrow the scope to the one class it's good at.
- [ ] Pick a threshold that matches the buyer's need (parental-control apps want high recall; brand-safety wants low false positives)

Phase 2 — Harden (2–3 lucid days)

- [ ] Verify the TFLite/quantized model matches the full model's accuracy (quantization can silently degrade it)
- [ ] Benchmark real latency on a target device; confirm the 30–50ms claim
- [ ] Robust preprocessing + graceful handling of weird inputs (corrupt files, odd sizes, grayscale)
- [ ] Wire in the OCR layer as an optional pass (text-in-image detection = a real differentiator)

Phase 3 — Productize (2–4 days)

- [ ] Choose the form factor (recommended: **hosted API first** — recurring revenue):
  - **Hosted API** → POST an image, get back `{safe: bool, scores: {...}}`. Charge per 1,000 calls or monthly tier. *Best recurring income.*
  - **Model + SDK license** → one-time sale (this is your ₹1,999 Instamojo listing — keep it running in parallel)
  - **Mobile SDK** → for on-device app builders (higher price, more support)
- [ ] Wrap the model in a small FastAPI service (I can scaffold this fully)
- [ ] Build a simple demo page — upload an image, see it flagged. **This is where a tool like base44 fits: the demo/dashboard shell.**
- [ ] Deploy it somewhere cheap (Render/Railway free tier to start)

Phase 4 — Sell it (ongoing)

- [ ] Name the buyer: parental-control apps, kids' platforms, dating apps, community/forum tools, image marketplaces, small app studios needing app-store compliance
- [ ] Pricing draft: Free tier (500 calls) → ₹ / $ per-call tiers → custom for volume
- [ ] Lead with the **accuracy number from Phase 1** + the **privacy angle** ("on-device / your images never leave your server")
- [ ] First outreach: 5–10 indie app devs who post about moderation pain

---

Who does what

The parts only YOU can do	The parts I can do for you
Source honest real-world test images	Build the eval harness + metrics script
Judge edge cases + set thresholds	Scaffold the FastAPI service + demo
Decide the buyer + the price	Write the README + docs
The outreach + the relationships	Clean/consolidate the codebase
Stand behind the accuracy claim	Draft benchmark + pricing pages

Honest risks & kill-criteria

- **If Phase 1 shows the model is weak and can't be fixed with reasonable data** → stop, pivot to selling it as the as-is ₹1,999 artifact, and put energy elsewhere. Validation *protects* you from pouring weeks into something unproven.
- **Don't let this eat the food clock.** Keep the fast track alive (Convergence + reels) while you build this over ~2 weeks of lucid windows.
- **This is a marathon lane, not the first-rupee lane.** First rupee still comes from reels/Convergence. SafeScreen is the *durable* income you build behind it.