GODSHAND — Cinematic AI Film Kit (₹499 Listing + Setup Sheet)

Turn a single portrait into a cinematic AI film shot — on a free GPU. GODSHAND is the exact pipeline used to render *The Saga of the Beholder*. You bring a character image; the kit gives you graded, HD, motion-consistent film frames without a render farm.

> Draft listing copy — rename "GODSHAND" freely. Everything below is written to be pasted into Topmate / Razorpay / Gumroad. The second half is the honesty sheet buyers should read *before* paying.

The one-liner

GODSHAND Cinematic AI Film Kit — a guided notebook that renders film-grade AI video shots from your own character images, running entirely on free Kaggle dual-T4 GPUs. No render farm, no paid cloud, no melt.

What it does

- Takes your **anchor image(s)** (a face/character) and animates them into short cinematic shots.
- Runs a hardened **Wan 2.2 image-to-video** pipeline (with a working LTXV fallback engine).
- Applies a full **film grade** — Kodachrome-style color, halation, film grain, sharpening — automatically.
- Upscales to **HD (1536×1024)** and interpolates for smooth motion.

- Renders in **shot batches** so you can build a whole film across multiple free-GPU sessions.

Why it's different (the hard-won part)

This isn't a fresh ComfyUI workflow you have to debug. It's a kit where the painful failures are **already solved**:

- ✅ **Anti-melt / anti-rainbow fix** — the render-schedule desync that turns AI video into noise soup is patched.
- ✅ **Memory-safe streaming grade** — the postprocess step no longer OOM-crashes on HD footage (bounded RAM, tunable).
- ✅ **Fits free 14 GB T4 GPUs** — quantized models chosen so it actually runs where you don't pay.
- ✅ **Shot/chunk system** — render 435-shot films in resumable batches.

What's inside

- The full `rainman-wan.ipynb` Kaggle notebook (drop-in, guided).
- The GODSHAND source kit (render, grade, upscale, face-enhance, stitch).
- A quick-start guide + the recommended settings that actually work.
- **A full video walkthrough** — the entire process screen-recorded end to end, so nothing is left to guesswork.
- Sample anchor images and a demo shot list.

Proof

The film speaks for the kit. *The Saga of the Beholder* — every frame rendered on free GPUs with this exact pipeline. **The demo clips are the use-case: film entertains, the kit empowers.**

Price

₹499 — one-time, lifetime access. Includes the full kit, the notebook, the quick-start guide, **and a complete step-by-step video walkthrough** of the entire

process — recorded end to end so you just follow along, click for click.

⚠️ Read before you buy — honest setup requirements

> 🧭 This is a **guided notebook**, not a one-click app. If you follow steps and have a little patience, it works. If you expect a magic "make my movie" button, this isn't that — and I'd rather you know now.

You will need

- ☐ A free **Kaggle account**, phone-verified (required to unlock GPUs).
- ☐ Kaggle's free **dual-T4 GPU** quota (~30 hrs/week — enough for real work, not unlimited).
- ☐ A free **HuggingFace account + access token** (to download the open models).
- ☐ Your own **character/anchor images** (or use the samples included).
- ☐ Basic comfort with: opening a notebook, pasting a token, clicking Run. **No coding required.**

Time & expectations

- Rendering is **slow on free GPUs** — roughly **8–9 minutes per shot**. A full film is many sessions, not one afternoon.
- Output is **stylized AI video**. Expect occasional imperfections (hands, tiny details, rare artifacts). The kit minimizes them; it doesn't make them impossible.

What this kit does NOT do

- ❌ It's not a hosted service — you run it on your own free Kaggle account.
- ❌ It won't run well on random low-VRAM setups; it's tuned for **Kaggle T4s**.
- ❌ It doesn't include a commercial-use guarantee for the underlying open models — **check each model's license** if you plan to sell what you render.

- ❌ It's not a chatbot or a script writer — it renders shots from images you provide.

FAQ

Do I need a registered business or GST to buy?

No — it's a digital download. You just need a way to pay.

Can I use my own faces/characters?

Yes. Drop in your own anchor images. Sample images are included to start.

Will it work on Google Colab?

It's built and tested for Kaggle dual-T4. Colab may work with tweaks but isn't supported out of the box.

Can I sell videos I make with it?

The kit code is yours to use, but the underlying AI models have their own licenses. Verify commercial terms for the models before selling output.

What support do I get?

Best-effort help via the setup guide + a contact channel. This is a solo-built kit, not a company helpdesk — priced accordingly.

Refund policy (suggested)

Because this is an instant digital download **with a complete video walkthrough included**, all sales are final once the files are accessed. There's no "I couldn't set it up" gap — the video shows every single step. If you still get stuck, reach out and I'll help you get running.

Pre-launch checklist

- [ ] Finish + export the film (your #1 proof asset).
- [ ] Pick a payment rail (Topmate / Razorpay individual / Gumroad).
- [ ] Record a 60–90s screen demo of the kit rendering one shot.
- [ ] Zip the kit + guide; test the download on a clean account.

- [ ] Publish listing → share proof clips → get first paying stranger.

🎥 Video walkthrough — recording script

> 🎬 Record on a **clean, fresh Kaggle account** (new login, no pre-setup). The buyer must see the exact path they'll walk — not your already-configured machine. Timelapse the slow render parts. Aim for **12–18 min** with chapter markers.

0. Cold open (10–20s) — the hook

Open on a finished graded shot from the film. Voiceover: "This was rendered on a *free* GPU, no paid cloud. In the next few minutes I'll show you exactly how — step by step, nothing skipped."

1. What you need (30s)

On screen: a browser, a free Kaggle account, a free HuggingFace account. "No installs. No coding. If you can paste a code and click Run, you can do this."

2. Kaggle account + GPU unlock (1–2 min)

Sign in → phone-verify → show where to enable the **dual-T4 GPU** accelerator. Call out the ~30 hrs/week free quota honestly.

3. HuggingFace token (1 min)

Create account → Settings → Access Tokens → create a read token → copy it. "This lets the notebook download the models."

4. Load the notebook + kit (1–2 min)

Import `rainman-wan.ipynb` → attach the GODSHAND kit dataset → add the token as a **Kaggle Secret** named `HF_TOKEN`. Show exactly where each field goes.

5. Add your character + story (2–3 min) — the magic moment

Upload your **anchor image(s)** → open the shot-list / story file → edit one shot's prompt so viewers see *their* story is what drives the render. This is the "it's *my* film" hook.

6. Run it (1 min + timelapse)

Click Run All. Narrate what's happening: models downloading, both experts loading, sampling starting. Timelapse the wait so it never feels slow.

7. First shot renders (1–2 min)

Show the raw base shot. Then the **HD upscale** and the **film grade** — cut a quick before/after so the color + grain "wow" lands.

8. Export + download (1 min)

Show the final `.mp4` and where to download it. "That's one shot. Repeat for your whole film — the kit resumes across sessions."

9. Outro (30s)

Recap the 3 things they needed. Point to the quick-start guide + how to get help. End on the finished film clip again. "That's the whole thing. Go make yours."

Recording tips

- Screen-record at 1080p; a small webcam bubble builds trust but is optional.
- Cut dead air ruthlessly; every model-download wait gets timelapsed.
- Add chapter markers matching the steps above — buyers rewatch specific parts.
- A 60–90s cut of steps 0, 5, 7 = your **sales demo**. Film both in one sitting.

🔴 How to record it (free)

Screen recorder: OBS Studio (free, all platforms) — or zero-install built-ins: Windows **Win+G**, Mac **Cmd+Shift+5**. Record at 1080p.

Don't film the slow render live. Record the setup steps → **stop** during the ~8–9 min render → **resume** for the result. Splice the segments (or drop a "rendering..." timelapse card) in editing.

Audio: narrate live with any headset/phone mic, *or* record silent and add captions + voiceover after — less pressure on take one.

Editing (free): CapCut, DaVinci Resolve, or iMovie / Clipchamp. **Import the actual film `.mp4` clips** for beats 0 and 7 — don't screen-record them, or they'll look soft.

Before you hit record: use a **clean browser profile** (bookmarks hidden, no personal tabs) and the **fresh Kaggle account**, so viewers see the real from-

scratch path.

🎯 Where to find buyers

🧭 Honest targeting note: the kit still needs a Kaggle account + HF token, so the real sweet spot isn't *literally* zero-tech — it's the **budget AI-creator who'll happily follow a video**. The faceless-YouTube / AI-video crowd is perfect: hungry, tutorial-native, allergic to paying for expensive tools.

💰 Demand is real: creators already pay for faceless-video workflows (one builder even got an $8K offer to buy their pipeline). Your edge is a one-time ₹499 kit that runs free on Kaggle — it directly beats the monthly SaaS tools this crowd resents.

Communities to target

Platform	Community	Why / rules to respect
Discord	Faceless Incubator (~38k)	Largest faceless-creator hub — biggest pool. Watch promo rules.
Discord	YouTube Faceless (~20.7k)	The #1 faceless community, but **strictly no self-promo** — value first, DM only if invited.
Discord	AI Forum for India (~7.5k)	India-focused AI hub — your budget-creator base.
Discord	AI Filmmakers server (invite pinned in r/aifilmmaking)	Craft-focused, platform-independent; share the film as proof.
Discord	Wan (official) / ComfyUI servers	The exact models your kit uses; showcase channels welcome results.
Reddit	r/aitubers	Small AI-video creators, tutorial-hungry. Showcase-friendly.
Reddit	r/aifilmmaking	Exactly your craft; people asking “how was this made?”

Platform	Community	Why / rules to respect
Reddit	r/comfyui, r/StableDiffusion	Technical AI-video crowd; best for "built on free GPUs" proof clips. No paywalled-workflow spam.
Reddit	r/FacelessAICreators, r/Seedance_AI	More faceless / AI-video subs; showcase-friendly.
India	AI Learners India (Hindi, no-code Discord)	Zero-tech-background, limited-funds beginners — your ideal buyer.
India	AI Video Generators (Telegram)	Active AI-video channel comparing tools; warm audience.
Facebook	"Faceless YouTube" creator groups	Older/less-technical creators on tight budgets.

How to find more (server-finder sites)

Search these by tag (`comfyui`, `filmmaking`, `ai-video`, `faceless`) to surface fresh servers: **Disboard**, **Discadia**, **top.gg**, **DigitalOcean's AI-Discord roundup**, and **aixploria's AI-Discord list**.

The value-first playbook (don't get banned)

1. **Lead with the film, not the link.** Post a stunning clip: "Rendered this entirely on a free Kaggle GPU." Let people ask *how*.
2. **Answer generously in public.** When someone asks about setup, actually help. You become "the free-GPU AI-film person."
3. **Soft-mention, don't spam.** "I packaged my exact pipeline + a full walkthrough video if anyone wants the shortcut — happy to DM."
4. **Use showcase channels** where self-promo *is* allowed; respect the no-promo rule everywhere else.
5. **Let the walkthrough clip do the selling** — people who see the whole process work are pre-sold and won't refund.

Ready-to-post templates

Reddit showcase post (r/aifilmmaking, r/StableDiffusion, r/aitubers)

> **Title:** I rendered a short film entirely on a *free* GPU — no paid cloud, no render farm

> Body: I spent months fighting the usual AI-video curse — melting faces, rainbow noise, characters that change identity every frame. I finally got a pipeline that holds a consistent face *and* grades like real film (Kodachrome warmth, grain, halation), running entirely on free Kaggle T4 GPUs.
> Here's a ~30-second piece from a longer memoir project I'm building.
> Ask me anything about how it's built — happy to share the setup. What would you want to see it render next?

(Attach the clip. No link in the post body — let people ask. Reply to "how?" with the soft-mention below.)

Discord showcase-channel message

> Made this on a **free** Kaggle GPU 👇 Took forever to kill the melt + rainbow noise, but it finally holds a real face and a proper film grade. Everything's open models on a free T4 — happy to answer setup questions 🙌

(Drop the clip. Only post in channels that allow showcases.)

Soft-mention reply (when someone asks "how did you do this?")

> Sure — it's Wan 2.2 image-to-video on ComfyUI, quantized to fit a free T4, with a custom grade + upscale pass. I packaged my exact notebook + a full step-by-step walkthrough video so people don't have to relive the debugging pain — happy to DM the link if that's useful, otherwise glad to explain any part here.